

Labeling in Strucplots

by David Meyer
Wirtschaftsuniversität Wien, Austria
David.Meyer@R-project.org

August 30, 2005

One of the major enhancements in **vcd** compared to `mosaicplot()` and `assocplot()` in base R is the labeling in the strucplot framework which now has become much more powerful and flexible. Like the panel, the shading, the spacing, and the legend, labeling is now completely modular. The user supplies either a labeling function, or, alternatively, a generating function that parameterizes a labeling function, to `strucplot()` which then draws the labels (generating functions are discussed in a separate vignette). The labeling is well-separated from the actual plotting that occurs in the panel functions; it actually only relies on the viewport tree produced by them, and the ‘dimnames’ attribute of the visualized table. The drawing of the labels happens after the actual plot has been drawn by the panel function. Thus, it is possible to supply one’s own labeling function, or to combine some of the basic functions to produce a more complex labeling.

In the following, we describe the three basic modules (`labeling_text()`, `labeling_list()`, and `labeling_cells()`) and derived functions that build upon them.

1 Labels in the borders: `labeling_text()`

`labeling_text()` is the default for all strucplot displays. It plots labels in the borders similar to the `mosaicplot()` function in base R, but is much more flexible: it is not limited to 4 dimensions, and the positioning and graphical parameters of levels and variable names are quite customizable. In addition, the problem of overlapping labels can be handled in several ways.

As an example, consider the ‘Titanic’ data, consisting of 4 categorical variables: Survived, Sex (i.e., ‘Gender’), Age, and Crew. By default, the variable names and levels are plotted ‘around’ the plot in a counter-clockwise way (see Figure 1):

```
> mosaic(Titanic)
```

Note that the last two levels of the ‘Survived’ variable do overlap, as well as some ‘Adult’ and ‘Child’ labels of the ‘Age’ Variable. This issue can be addressed in several ways. The ‘brute force’ method is to enable clipping for these dimensions (see Figure 2):

```
> mosaic(Titanic, labeling_args = list(clip = c(Survived = TRUE,  
+      Age = TRUE)))
```

The `clip` parameter is passed to the labeling function via the `labeling_args` argument which takes a list of parameters. `clip` itself takes a vector of logicals (one for each dimension). Almost all vectorized arguments in the strucplot framework can be abbreviated in the following way: unnamed components (or the defaults, if there are none) are recycled as needed, but overridden by the named components. Here, the default is `FALSE`, and therefore clipping is enabled only for the ‘Survived’ and ‘Age’ dimensions. A more sensible solution to the overlap problem is to abbreviate the levels (see Figure 3):

```
> mosaic(Titanic, labeling_args = list(abbreviate = c(Survived = TRUE,  
+      Age = 4)))
```

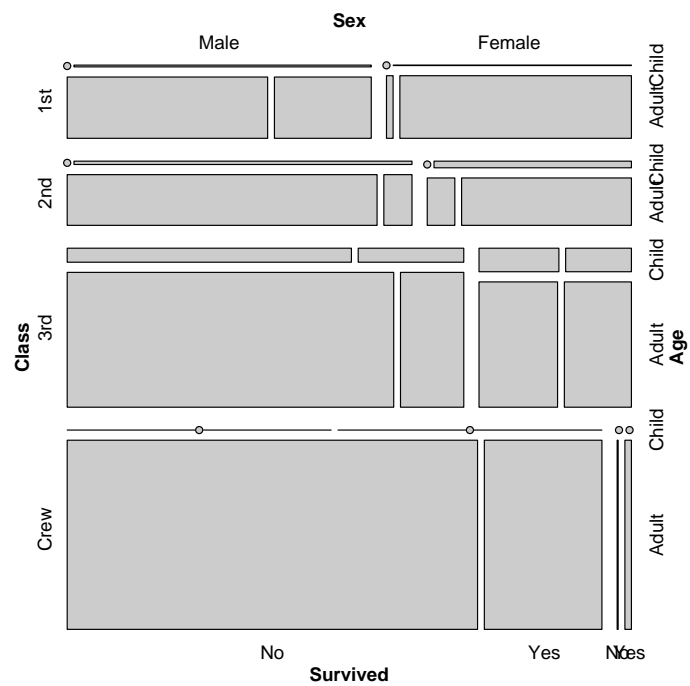


Figure 1: Mosaic plot for the Titanic data.




Figure 2: The effect of clipping.

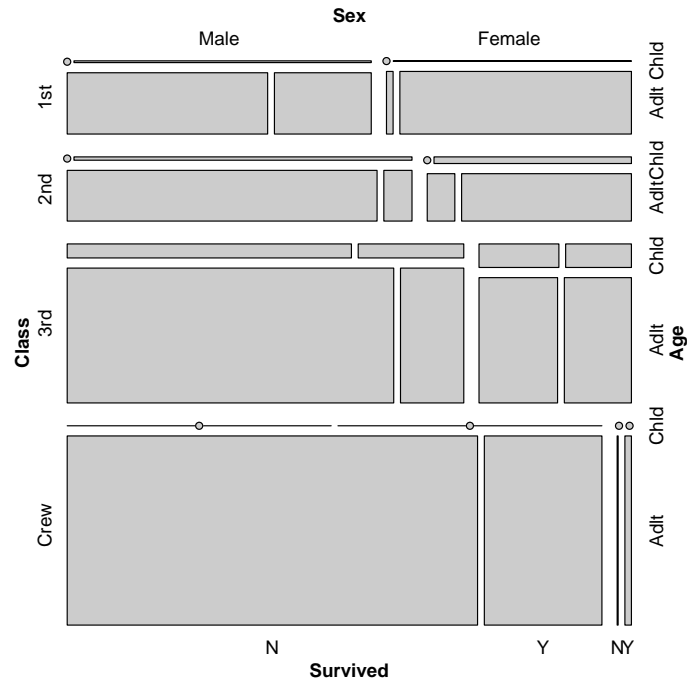


Figure 3: Abbreviating.

The `abbreviate` argument takes a vector of integers indicating the number of significant characters the levels should be abbreviated to (TRUE is interpreted as 1, obviously). Abbreviation is performed using the `abbreviate()` function in base R. Another possibility is to rotate the levels (see Figure 4):

```
> mosaic(Titanic, labeling_args = list(rot_labels = c(bottom = 90,
+   right = 0)), margins = c(right = 6, bottom = 4))
```

Finally, we could also inhibit the output of repeated levels (see Figure 5):

```
> mosaic(Titanic, labeling_args = list(rep = c(Survived = FALSE,
+   Age = FALSE)))
```

We now proceed with a few more ‘cosmetic’ features. A first simple, but effectful modification is to position all labels and variables left-aligned, which can also be achieved by using `labeling_left()` (see Figure 6):

```
> mosaic(Titanic, labeling_args = list(pos_varnames = "left",
+   pos_labels = "left", just_labels = "left", rep = FALSE))
```

Note that obviously we need to change the justification to "left" as well. Next, we show how to put all levels to the bottom and right margins, and all variable names to the top and left margins (see Figure 7):

```
> mosaic(Titanic, labeling_args = list(tl_labels = FALSE,
+   tl_varnames = TRUE, abbreviate = c(Survived = 1,
+   Age = 4)))
```

The `tl_foo` (“top left”) arguments are TRUE by default. Now, we will add boxes to the labels and additionally enable clipping (see Figure 8):



Figure 4: Rotating labels.



Figure 5: Inhibiting the repetition of levels.

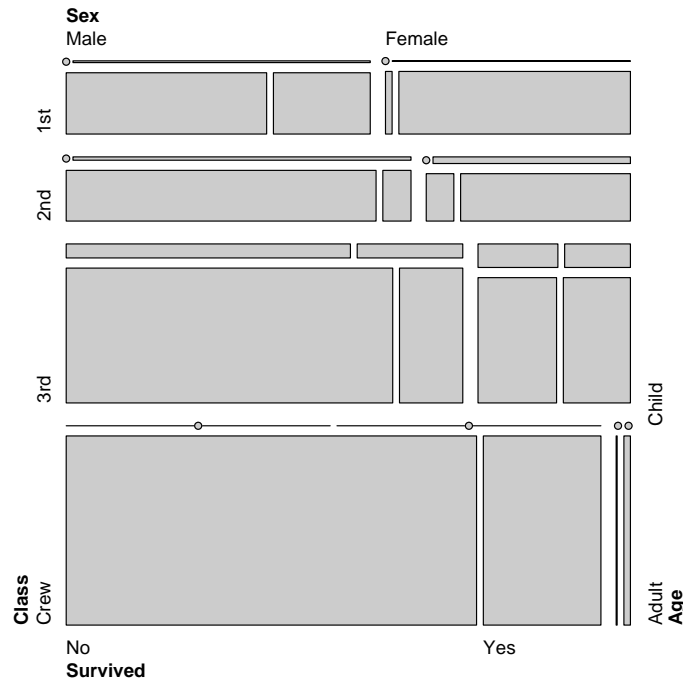


Figure 6: Left-aligning.

```
> mosaic(Titanic, labeling_args = list(tl_labels = FALSE,
+   tl_varnames = TRUE, boxes = TRUE, clip = TRUE))
```

The values to `boxes` and `clip` are recycled for all dimensions. The result is pretty close to what the `labels_cboxed()` wrapper does. Another variant is to put the variable names into the same line as the levels (see Figure 9):

```
> mosaic(Titanic, labeling_args = list(tl_labels = TRUE,
+   boxes = TRUE, clip = TRUE, labbl_varnames = TRUE),
+   margins = c(left = 4, right = 1, 3))
```

`labbl_varnames` (“variable names to the bottom/left of the labels”) is a vector of logicals indicating the side for the variable names. The resulting layout is basically what `labeling_lboxed()` produces, and a similar design is used by the `doubledecker()`.

2 Labels in the cells: `labeling_cells()`

This labeling draws both variable names and levels in the cells. As an example, we use the ‘PreSex’ data on pre- and extramarital sex and divorce (see Figure 10):

```
> mosaic(~MaritalStatus + Gender, data = PreSex, labeling = labeling_cells)
```

In the case of narrow cells, it might be useful to abbreviate labels and/or variable names and turn off clipping (see Figure 11):

```
> mosaic(~PremaritalSex + ExtramaritalSex, data = PreSex,
+   labeling = labeling_cells(abbreviate_labels = TRUE,
+   abbreviate_varnames = TRUE, clip = FALSE))
```




Figure 7: Changes in the margins.

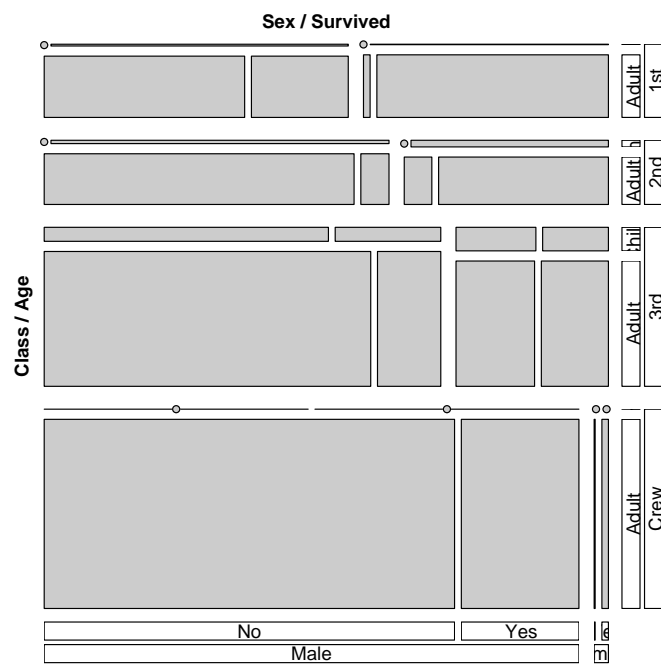


Figure 8: Boxes and Clipping.




Figure 9: Variable names beneath levels.

For some data, it might be convenient to combine cell labeling with border labeling as done by `labels_conditional()` (see Figure 12):

```
> mosaic(~PremaritalSex + ExtramaritalSex | MaritalStatus +
+   Gender, data = PreSex, labeling = labeling_conditional(abbreviate_varnames = TRUE,
+   abbreviate_labels = TRUE))
```

Additionally, the cell labeling allows the user to add arbitrary text to the cells by supplying a character array in the same shape than the data array to the `text` argument (cells with missing values are ignored). In the following example, this is used to add all observed values greater than 5 to the cells after the mosaic for the ‘Titanic’ data has been plotted (see Figure 13):

```
> mosaic(Titanic, pop = FALSE, labeling_args = list(abbreviate = c(Survived = 1,
+   Age = 4)))
> tab <- ifelse(Titanic < 6, NA, Titanic)
> labeling_cells(text = tab, clip = FALSE)(Titanic)
```

3 A simple list of labels: `labels_list()`

If problems with overlapping labels cannot satisfactorily resolved, the last remedy could be to simply list the levels below the plot (see Figure 14):

```
> mosaic(Titanic, labeling = labeling_list, margins = c(bottom = 5))
```

The number of columns can be specified.

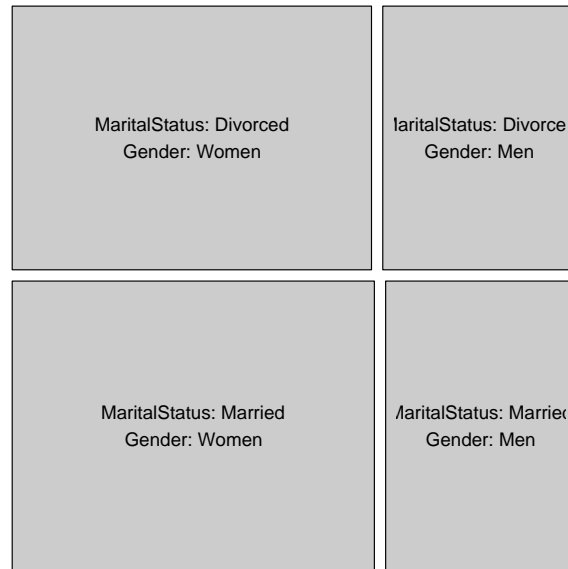


Figure 10: Cell labeling for the 'PreSex' data.

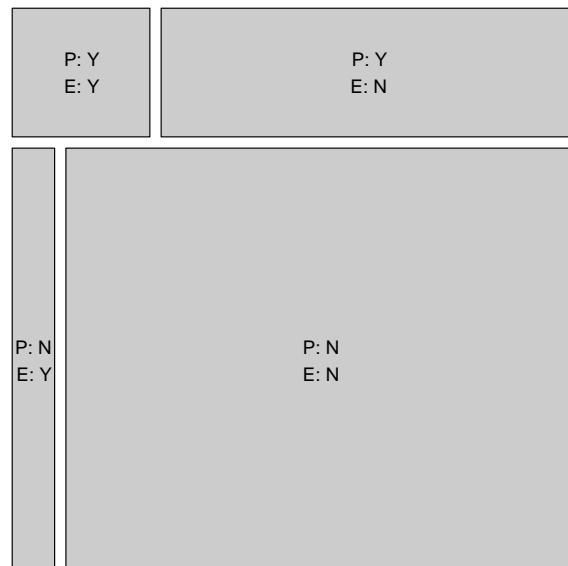


Figure 11: Cell labeling for the 'PreSex' data, labels abbreviated.

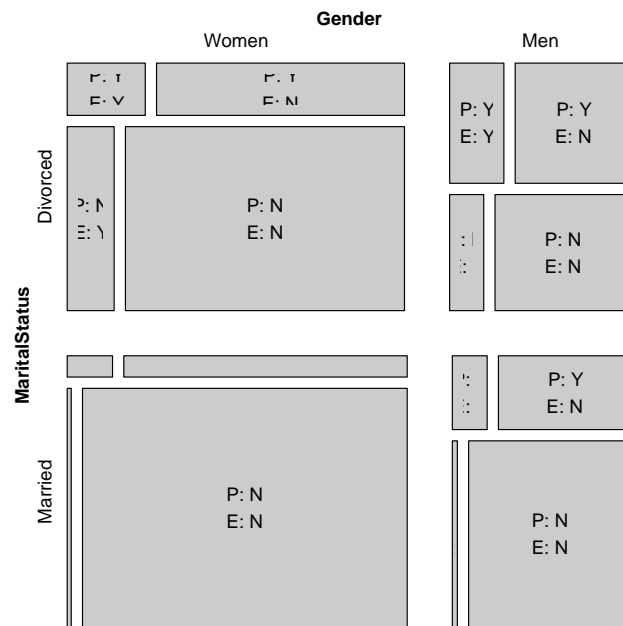


Figure 12: Conditional labeling for the 'PreSex' data, labels abbreviated.

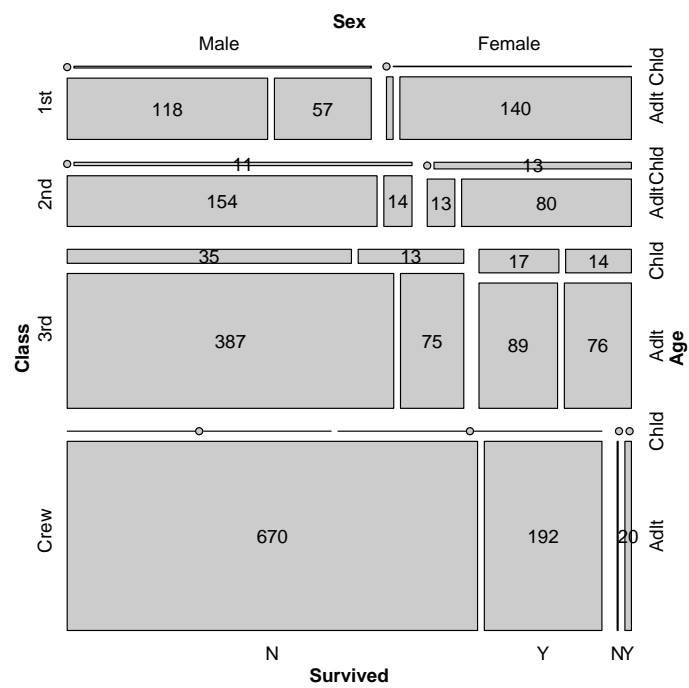


Figure 13: User-supplied Text added to a mosaic display of the 'Titanic' data.

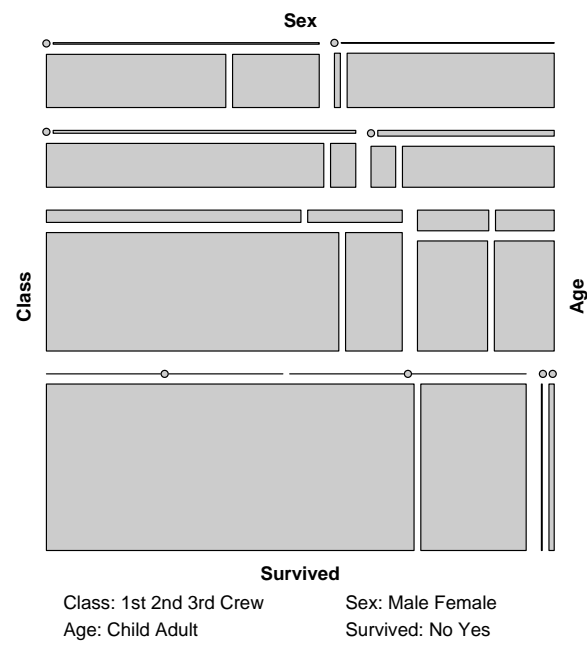


Figure 14: Labels indicated below the plot.